



# MapReduce and Spark: Overview

2015

Professor Sasu Tarkoma

# Overview

Computing Environment

History of Cluster Frameworks

Hadoop Ecosystem

Overview of State of the Art

MapReduce Explained

# Computing Environment

Scaling up

- More powerful servers

Scaling out

- More servers

Clusters provide computing resources

- Space requirements, power, cooling

- Most power converted into heat

Datacenters

- Massive computing units

- Warehouse-sized computer** with hundreds or thousands of racks

Networks of datacenters



# Cluster Computing Environment

Big Data compute and storage nodes are stored on racks based on common off the shelf components

Typically many racks in a cluster or datacenter

The compute nodes are connected by a high speed network (typically 10 Gbit/s Ethernet)

Different datacenter network topologies

Intra-rack and inter-rack communication have differing latencies

Nodes can fail

Redundancy for stored file (replication)

Computation is task based

**Software ensures fault-tolerance and availability**

# Typical Hardware

CSC Pouta Cluster running on the Taito supercluster in Kajaani

The nodes are HP ProLiant SL230s servers with two Intel Xeon 2.6 GHz E5-2670 CPUs

16 cores per server

Most with 64 GB of RAM per server

Taito extension in 2014: 17 000 cores

The nodes are connected using a fast FDR InfiniBand fabric

# Cloud Computing

Definition by NIST:

Cloud computing is a model for enabling convenient, on-demand network access to a shared pool of configurable computing resources (e.g., networks, servers, storage, applications, and services) that can be rapidly provisioned and released with minimal management effort or service provider interaction.

IaaS, PaaS, SaaS, XaaS

Big Data Frameworks are typically run in the cloud

# Big Data Environment

Typically common-of-the-shelf servers

Compute nodes, storage nodes, ...

Virtualized resources running on a cloud platform

Heterogeneous hardware, choice of OS

Contrasts traditional High Performance Computing (HPC)

# History of Cluster Frameworks

2003: Google GFS

2004: Google Map-Reduce

2005: Hadoop development starts

2008: Apache Hadoop (in production)

2008: Yahoo! Pig language

2009: Facebook Hive for data warehouses

2010: Cloudera Flume (message interceptor/filtering model)

2010: Cloudera S4 (continuous stream processing)

2011: LinkedIn Kafka (topic-based commit log service)

2011: Storm (Nathan Marz)

2011: Apache Mesos cluster management framework

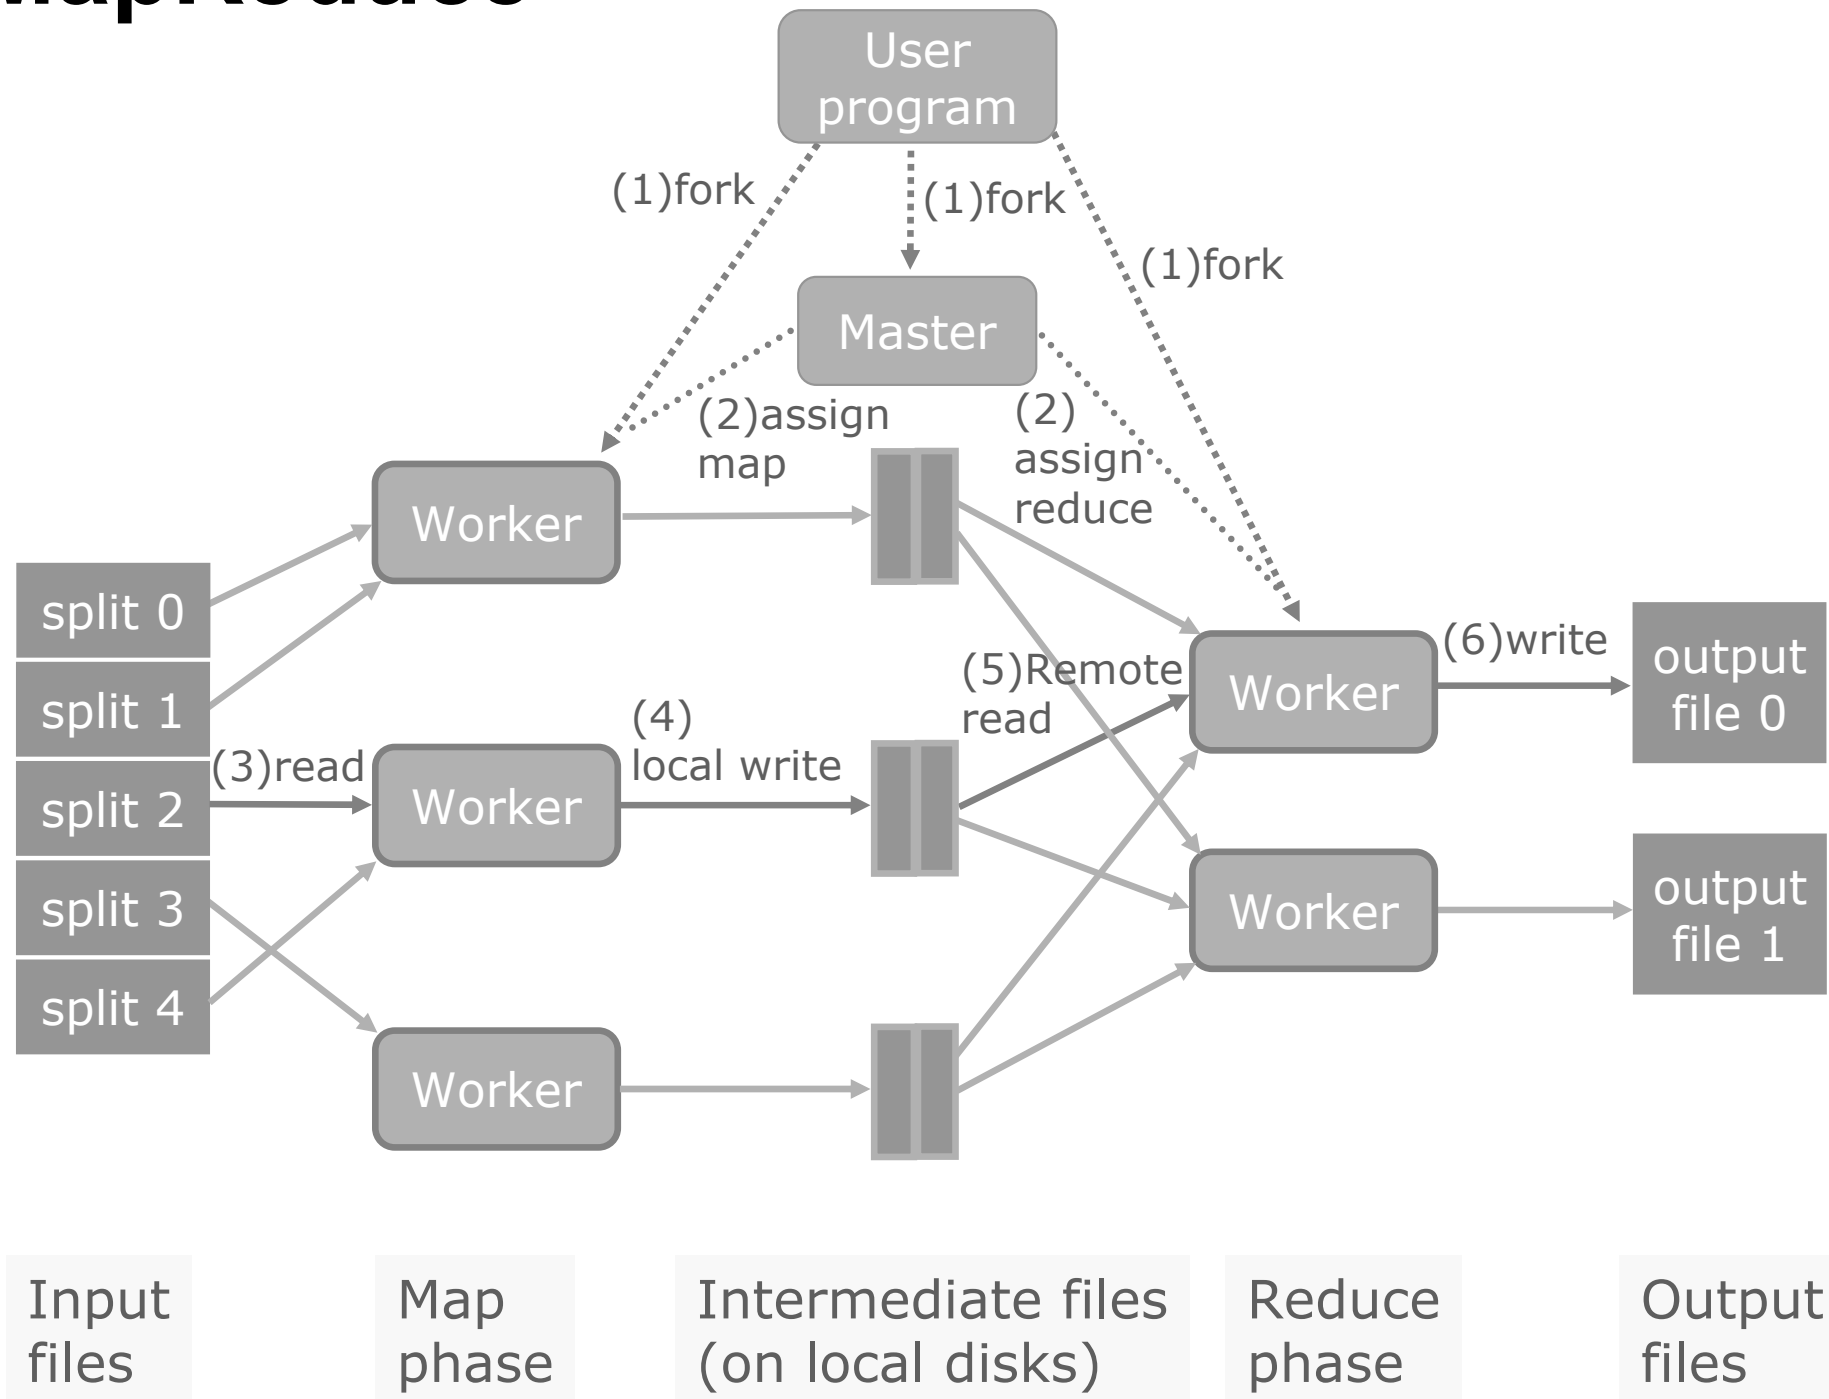
2012: Lambda Architecture (Nathan Marz)

2012: Spark for iterative cluster programming

2013: Shark for SQL data warehouses



# MapReduce



# **Major trends**

## **Apache Hadoop**

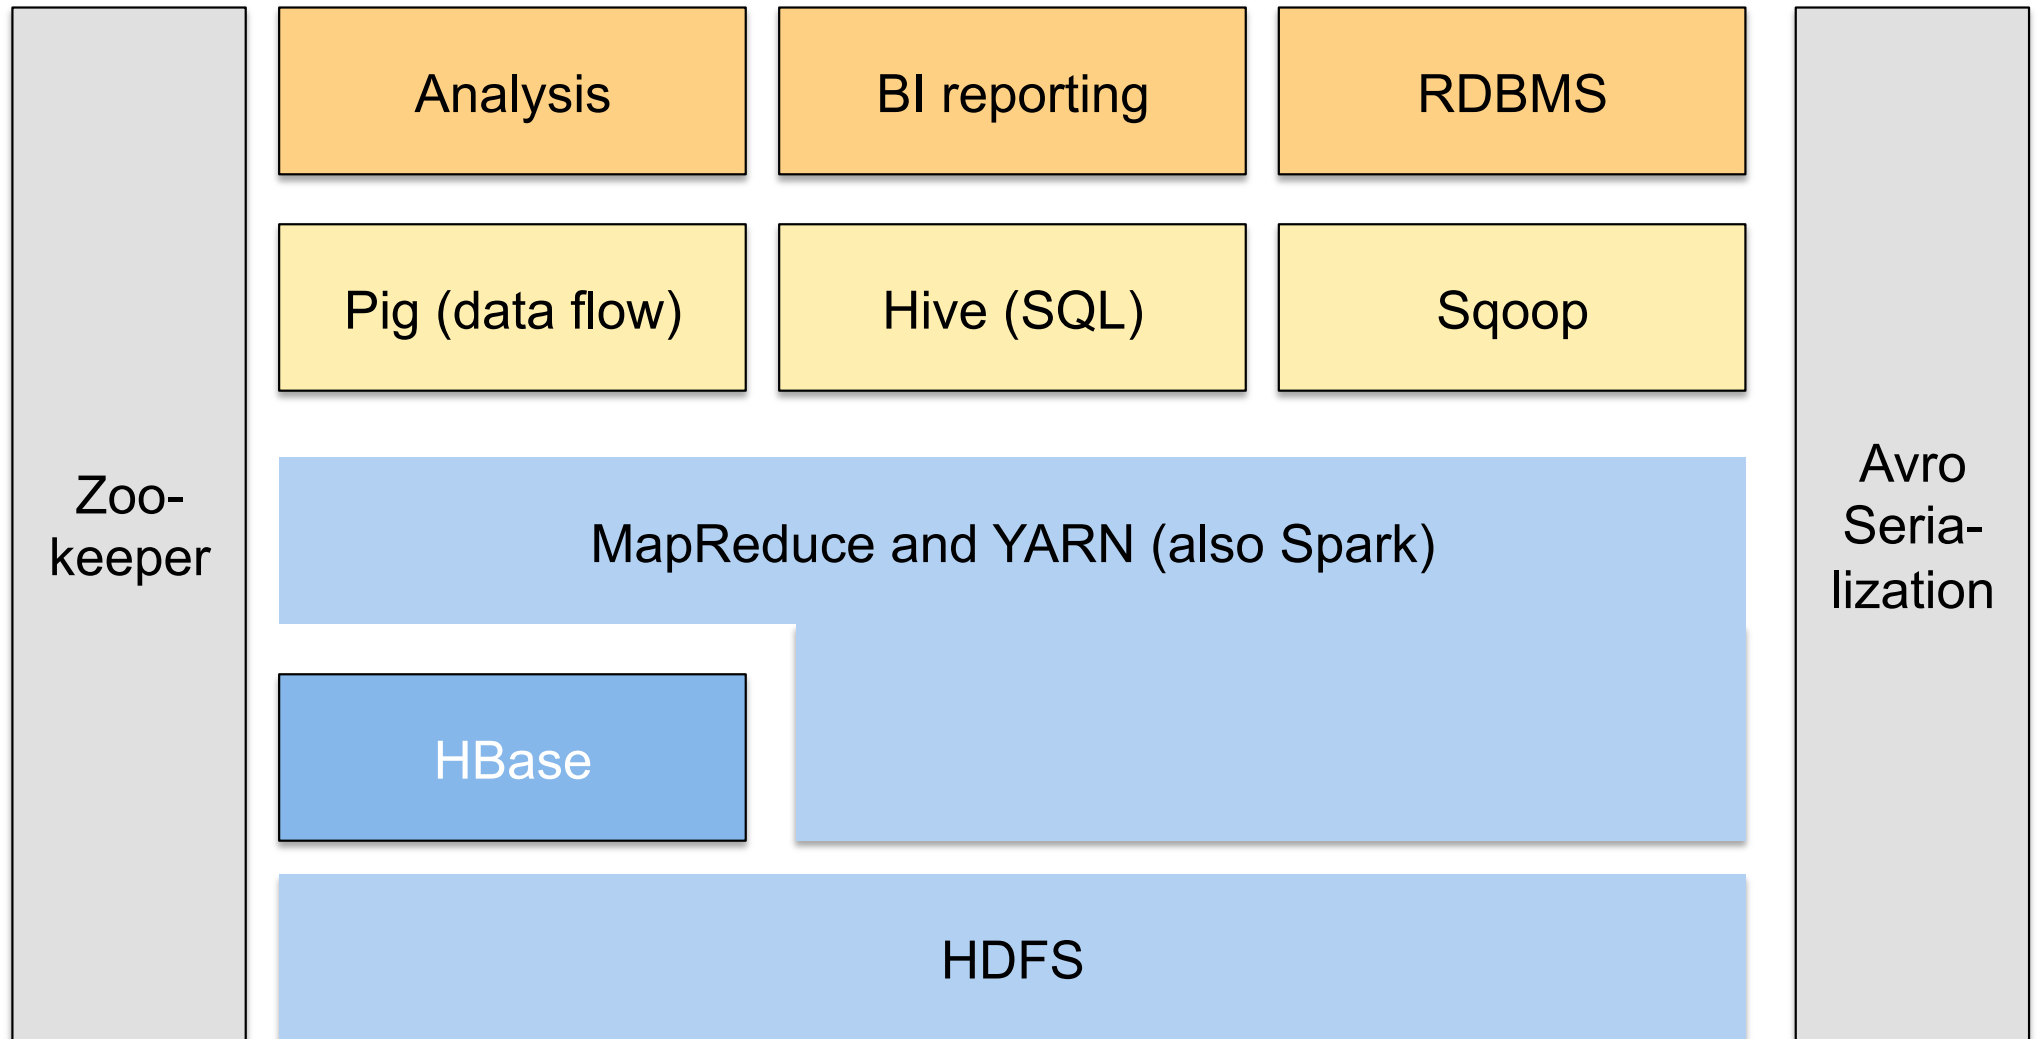
Hive, R, and others

## **Berkeley Data Analytics Stack (BDAS)**

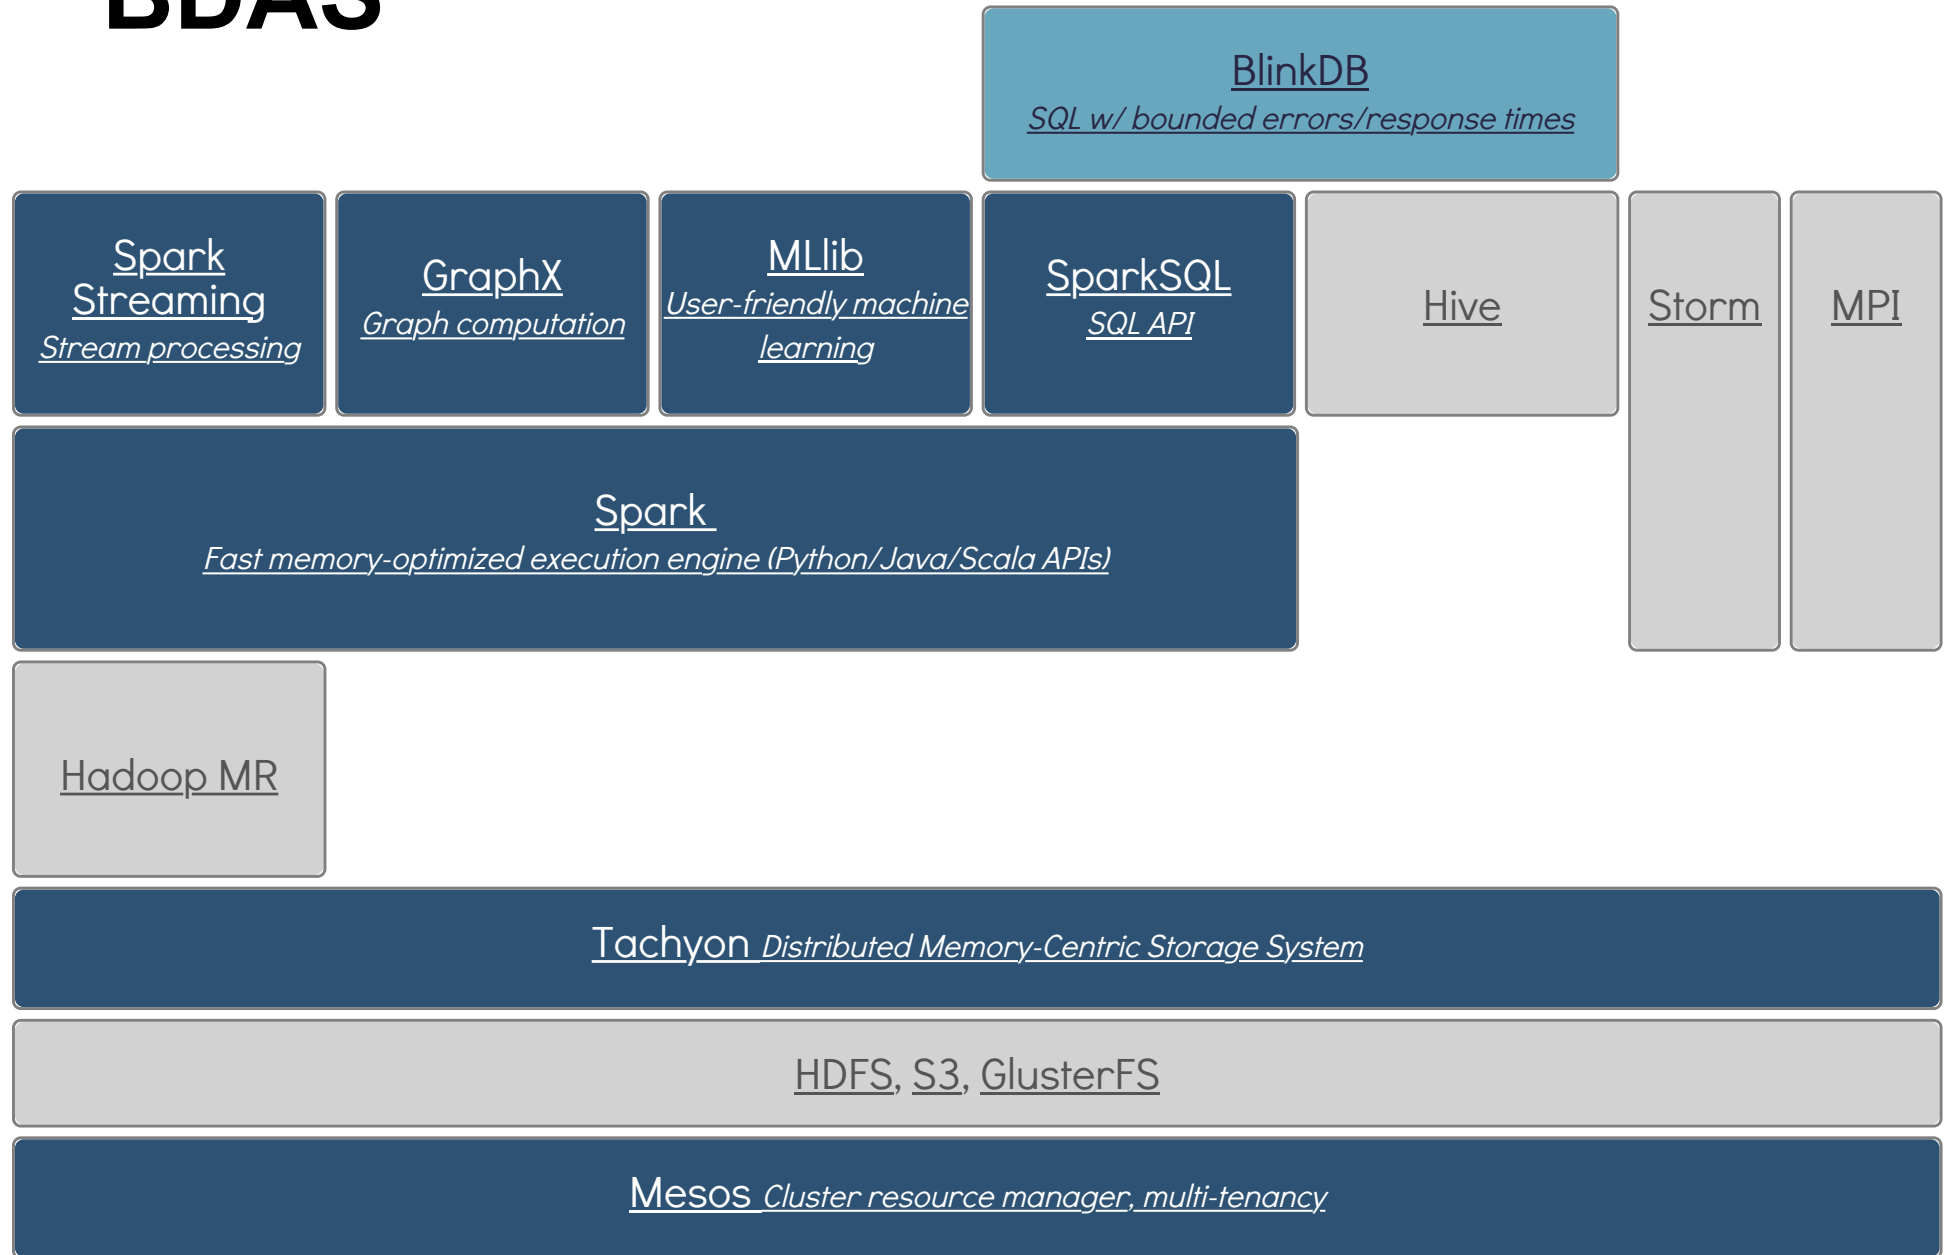
Mesos, Spark, Mlib, GraphX, Shark, ...

**Apache Spark is part of Apache Hadoop**

# Apache Hadoop Ecosystem



# BDAS



Supported Release



In Development



Related External Project

# Key idea in Spark

## **Resilient distributed datasets (RDDs)**

Immutable collections of objects across a cluster

Built with parallel transformations (map, filter, ...)

Automatically rebuilt when failure is detected

Allow persistence to be controlled (in-memory operation)

## **Transformations** on RDDs

Lazy operations to build RDDs from other RDDs

Always creates a new RDD

## **Actions** on RDDs

Count, collect, save

# Overview of State of the Art

- Data storage
- Data storage for real-time
- Data analysis
- Real-time data analysis
- Statistics and machine learning

# State of the Art: Data Storage

## **GFS (Google File System) and HDFS (Hadoop Distributed File System)**

Data replicated across nodes

HDFS: rack-aware placement (replicas in different racks)

Take data locality into account when assigning tasks

Do not support job locality (distance between map and reduce workers)

## **Hbase**

Modeled after Google's BigTable for sparse data

Non-relational distributed column-oriented database

Rows stored in sorted order

## **Sqoop**

Tool for transferring data between HDFS/Hbase and structural datastores

Connectors for MySQL, Oracle, ... and Java API

# Example: HDFS Architecture

HDFS has a master/slave architecture

NameNode is the master server for metadata

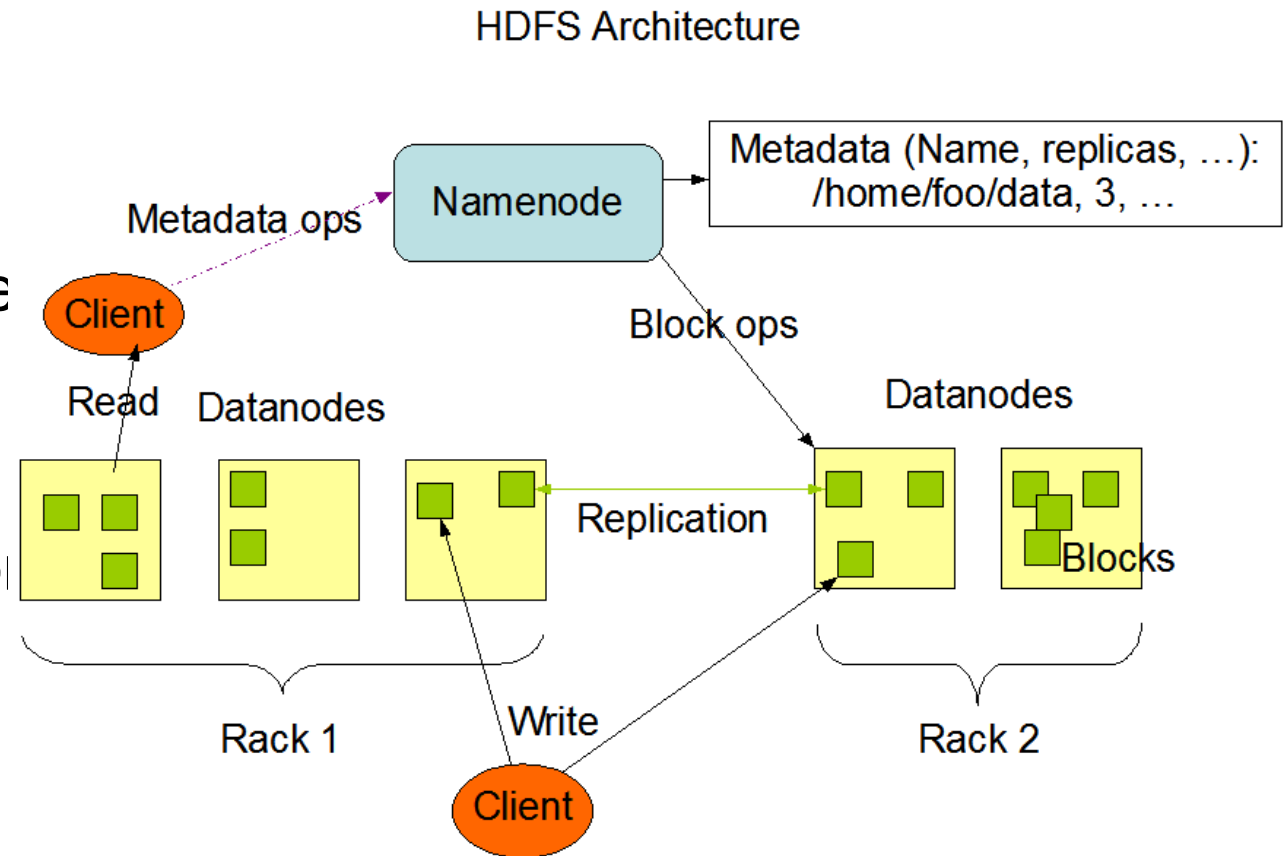
DataNodes manage storage

A file is stored as a sequence of blocks

The blocks are replicated for fault-tolerance

Common replication scheme: factor of 3, one replica local, two in a remote rack

**Rack-aware replica placement**



NameNode provides information for retrieving blocks  
Nearest replica is used to retrieve a block



# State of the Art: Data Storage for Real-time

## Kafka

Distributed, partitioned, replicated commit log service

Keeps messages in categories

Topic based system

Coordination through Zookeeper (through distributed consensus)

## Kestrel

Distributed message queue (server has a set of queues)

A server maintains queues (FIFO)

Does not support ordered consumption

Simpler than Kafka

# State of the Art: Data Analysis I/II

## MapReduce

Map and reduce tasks for processing large datasets in parallel

## Hive

A data warehouse system for Hadoop

Data summarization, ad-hoc queries, analysis for large sets

SQL-like language called HiveQL

## Pig

Data analysis platform

High-level language for defining data analysis programs, Pig Latin, procedural language

## Cascading

Data processing API and query planner for workflows

Supports complex Hadoop Map-Reduce workflows

## Apache Drill

SQL query engine for Hadoop and noSQL

# State of the Art: Data Analysis II

## Spark

- Cluster computing for data analytics

- In-memory storage for iterative processing

## Shark

- Data warehouse system (SQL) for Spark

- Up to 100x faster than Hive

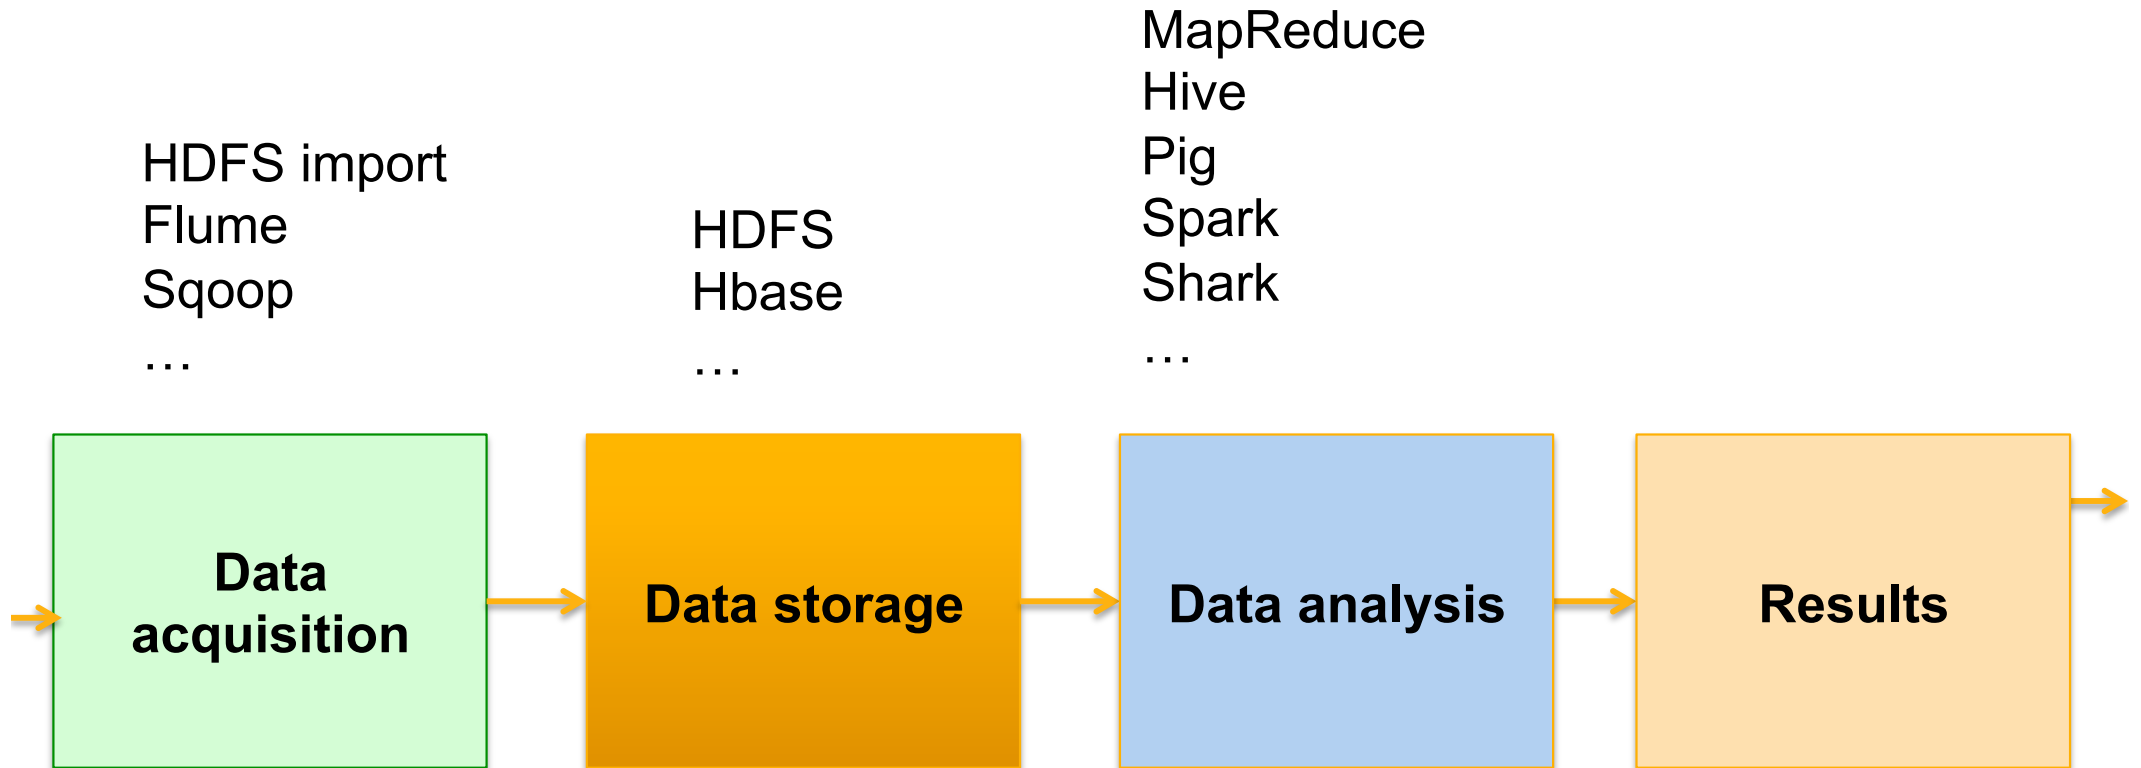
Spark/Shark is a distinct ecosystem from Hadoop

- Faster than Hadoop

- Support for Scala, Java, Python

- Can be problematic if reducer data does not fit into memory

# Summary of batch systems



# State of the Art: Real-time Data Analysis I/II

## Flume

Interceptor model that modifies/drops messages based on filters

Chaining of interceptors

Combine with Kafka

## Storm

Distributed realtime computation framework

“Hadoop for realtime”

Based on processing graph, links between nodes are streams

## Trident

Abstraction on top of Storm


Operations: joins, filters, projections, aggregations, ..

Exactly once-semantics (replay tuples for fault tolerance, stores additional state information)

<https://storm.apache.org/documentation/Trident-state>

# Flume example

Source	Channel	Sink
Avro	Memory	HDFS
Thrift	JDBC	Logger
Exec	File	Avro
HTTP		Null
JMS		Thrift
Syslog TCP/IP		File roll
		Hbase
Custom		Custom



# Storm

Developed around 2008-2009 at BackType, open sourced in 2011

**Spout:** is a flow of tuples

**Bolt:** accepts tuples and operates on those

**Topologies:** spouts → bolts → spouts

Example:

Tweet spout → parse Tweet bolt → count hashtags Bolt

Tweet spout → store in a file

# State of the Art: Real-time Data Analysis II

## Simple Scalable Streaming System (S4)

Platform for continuous processing of unbounded streams

Based on processing elements (PE) that act on events (key, attributes)

## Spark streaming

Spark for real-time streams

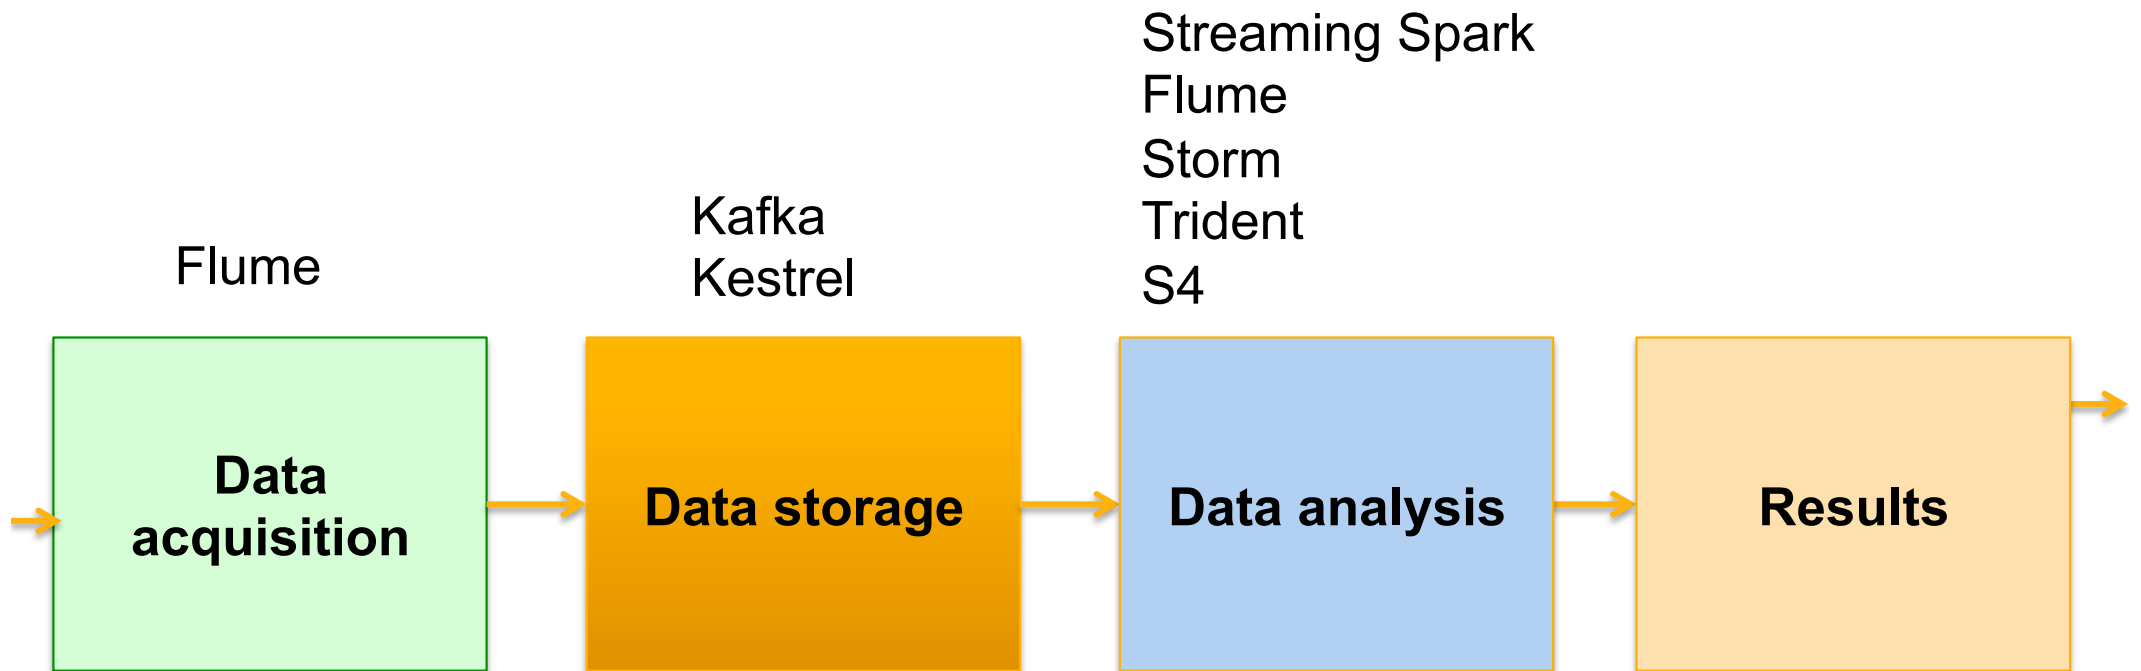
Computation with a series of short batch jobs (windows)

State is kept in memory

API similar to Spark



# Summary of real-time processing



# State of the art: Hybrid models

**Lambda architecture** combined batch and stream processing

Supports volume (batch) + velocity (streaming)

## Hybrid models

**SummingBird** (Hadoop + Storm)

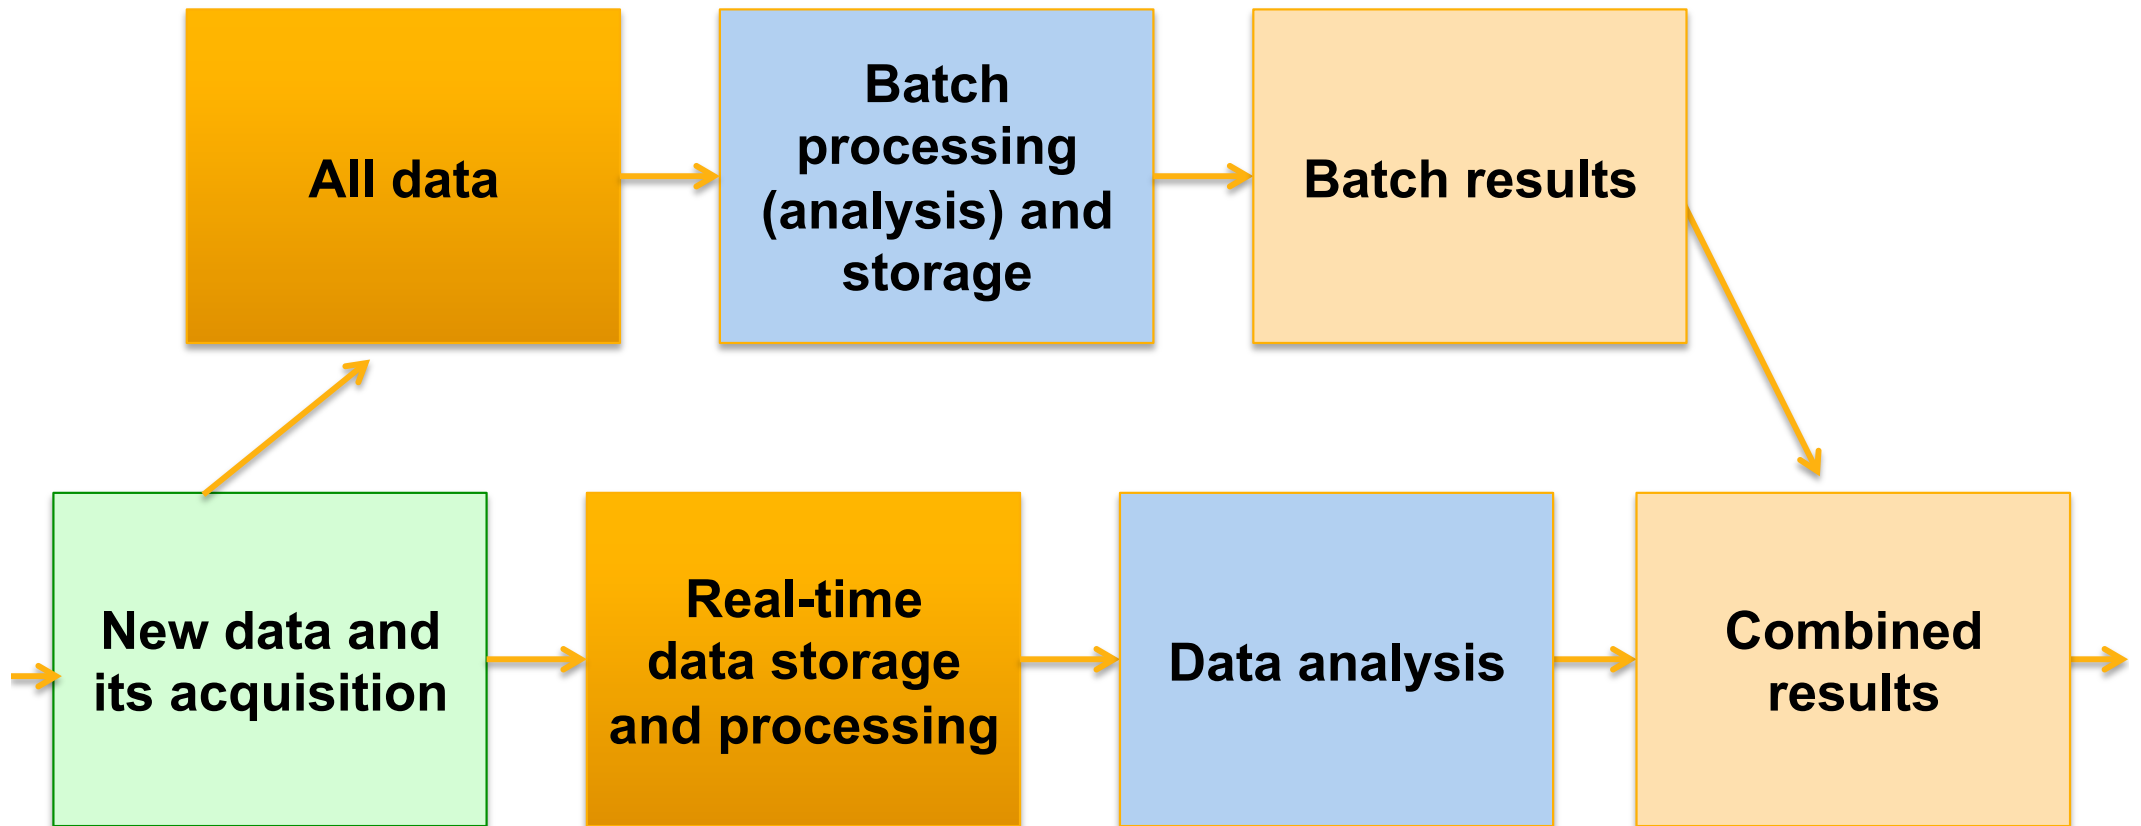
MapReduce like process with Scala syntax

**Lambdoop** (abstraction over Hadoop, HBase, Sqoop, Flume, Kafka, Storm, Trident)

Common patterns provided by platform

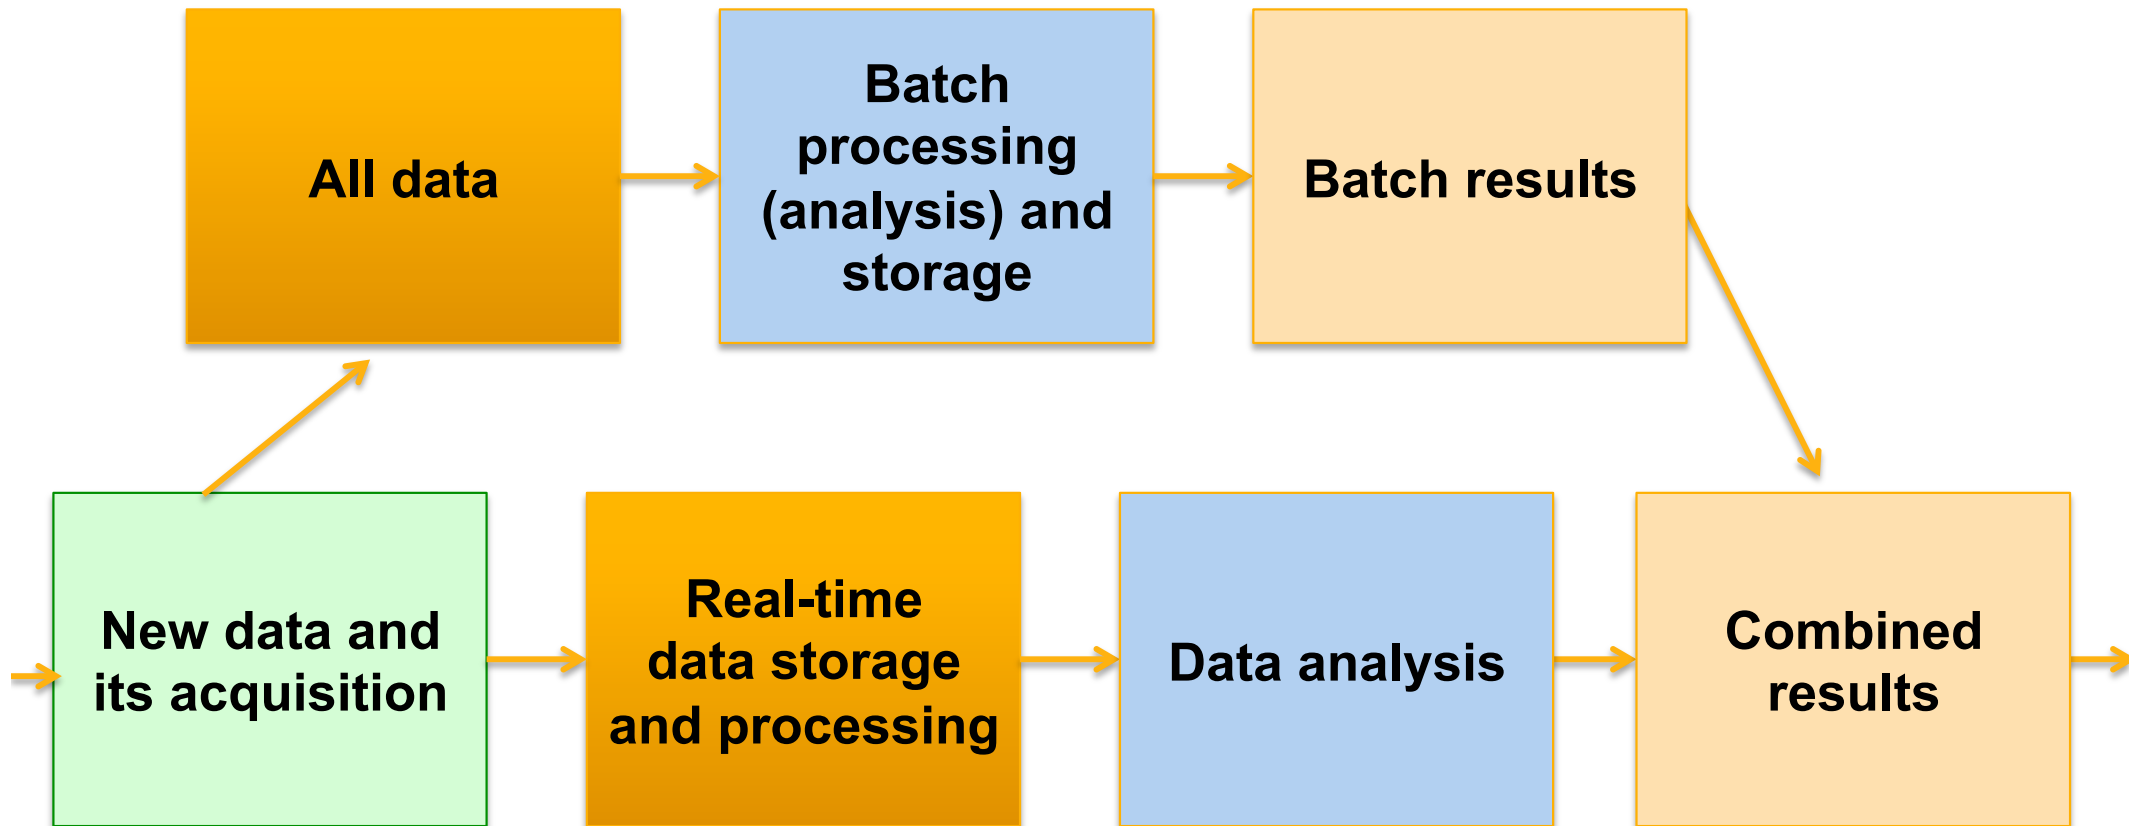
No MapReduce like process

# Lambda Architecture



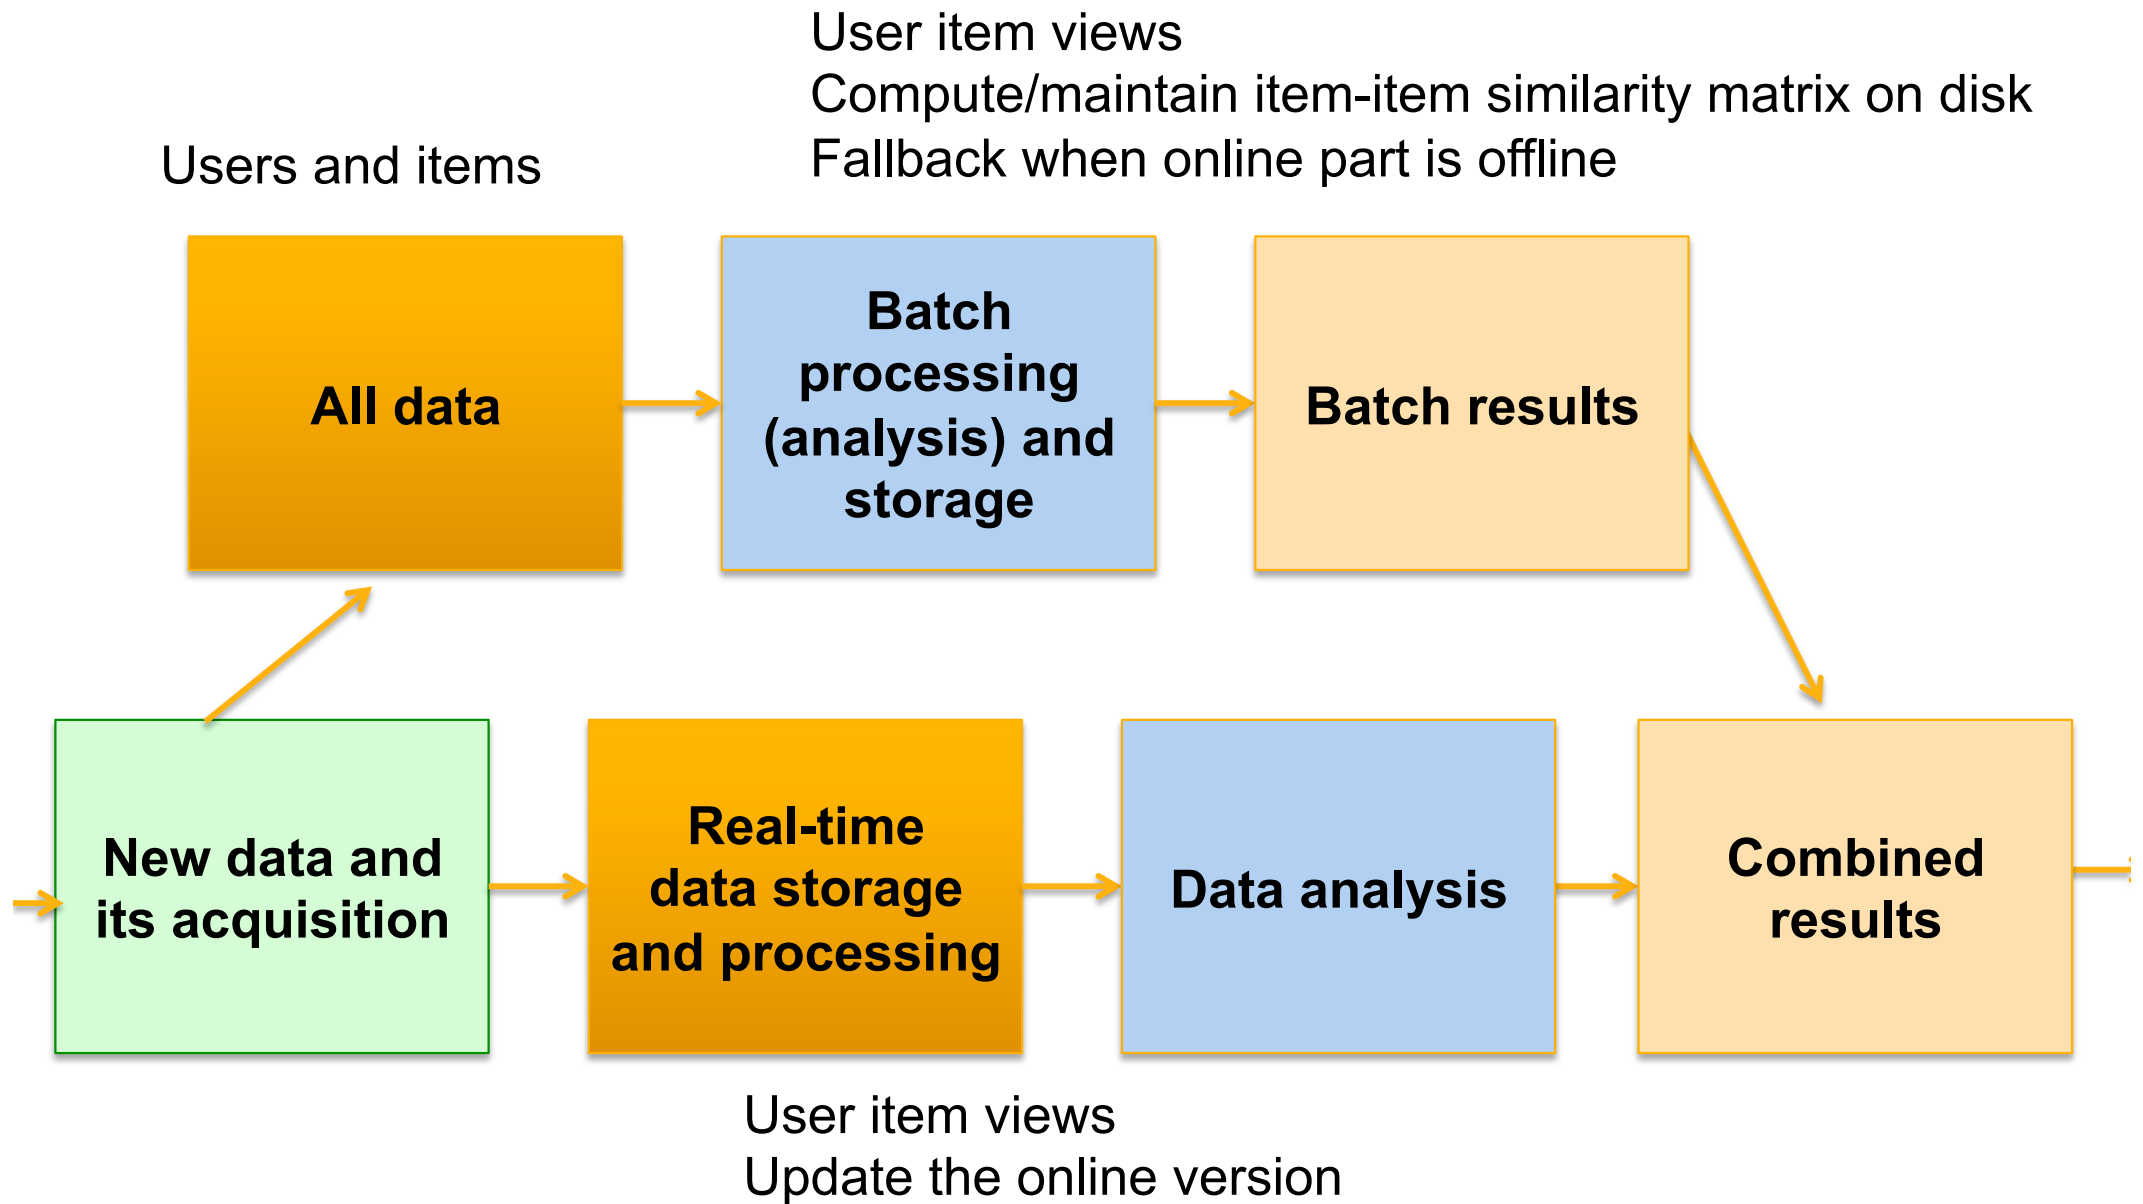
# Lambda Architecture: Twitter hashtags

Compute every hour the hashtag counts for the last hour stored on disk



Compute every five minutes the hashtag counts for the last five minutes stored in memory

# Lambda Architecture: Recommendations



# Challenges for the platform

Exactly-once semantics

Requires costly synchronization

High velocity: how to go to thousands of messages per second

Changes to structures and schemas

Data versioning in a production system

# State of the Art: Statistics and Machine Learning

## **R for Hadoop**

Distributed R for the cluster environment

## **R for Spark**

## **Mahout**

Currently Hadoop, next Spark

## **Weka**

State of the art machine learning library

Does not focus on the distributed case

Hadoop support, Spark wrappers

## **MLLib**

Machine learning for Spark

# State of the Art Distributed Toolbox

High-level applications

Hybrid systems (Hadoop+Storm, Spark + Spark streaming), optimization tier

Statistics and machine learning tier

Hive

R

Shark

Mlib

R

Trident

Task distribution tier

Hadoop, YARN

Spark  
Mesos

Storm

Spark  
Stream-  
ing

Storage tier

Storage  
GFS, HDFS, HBase

Real-time storage  
Kafka, Kestrel





# MapReduce

**2015**

**Professor Sasu Tarkoma**

# MapReduce Model

Google MapReduce introduced in 2004

Jeffrey Dean et al. MapReduce: Simplified Data Processing on Large Clusters. OSDI 2004.

Apache Hadoop since 2005

<http://hadoop.apache.org/>

Apache Hadoop 2.0 introduced in 2012

Vinod Kumar Vavilapalli et al. Apache Hadoop YARN: Yet Another Resource Negotiator, SOCC 2013.

New cluster resource management layer (YARN)

# Implicit Parallelism

The map function has implicit parallelism

This is because the order of the application of the function  $f$  to elements in a list is commutative

We can parallelize or reorder the execution

**MapReduce builds on this parallelism**

# MapReduce

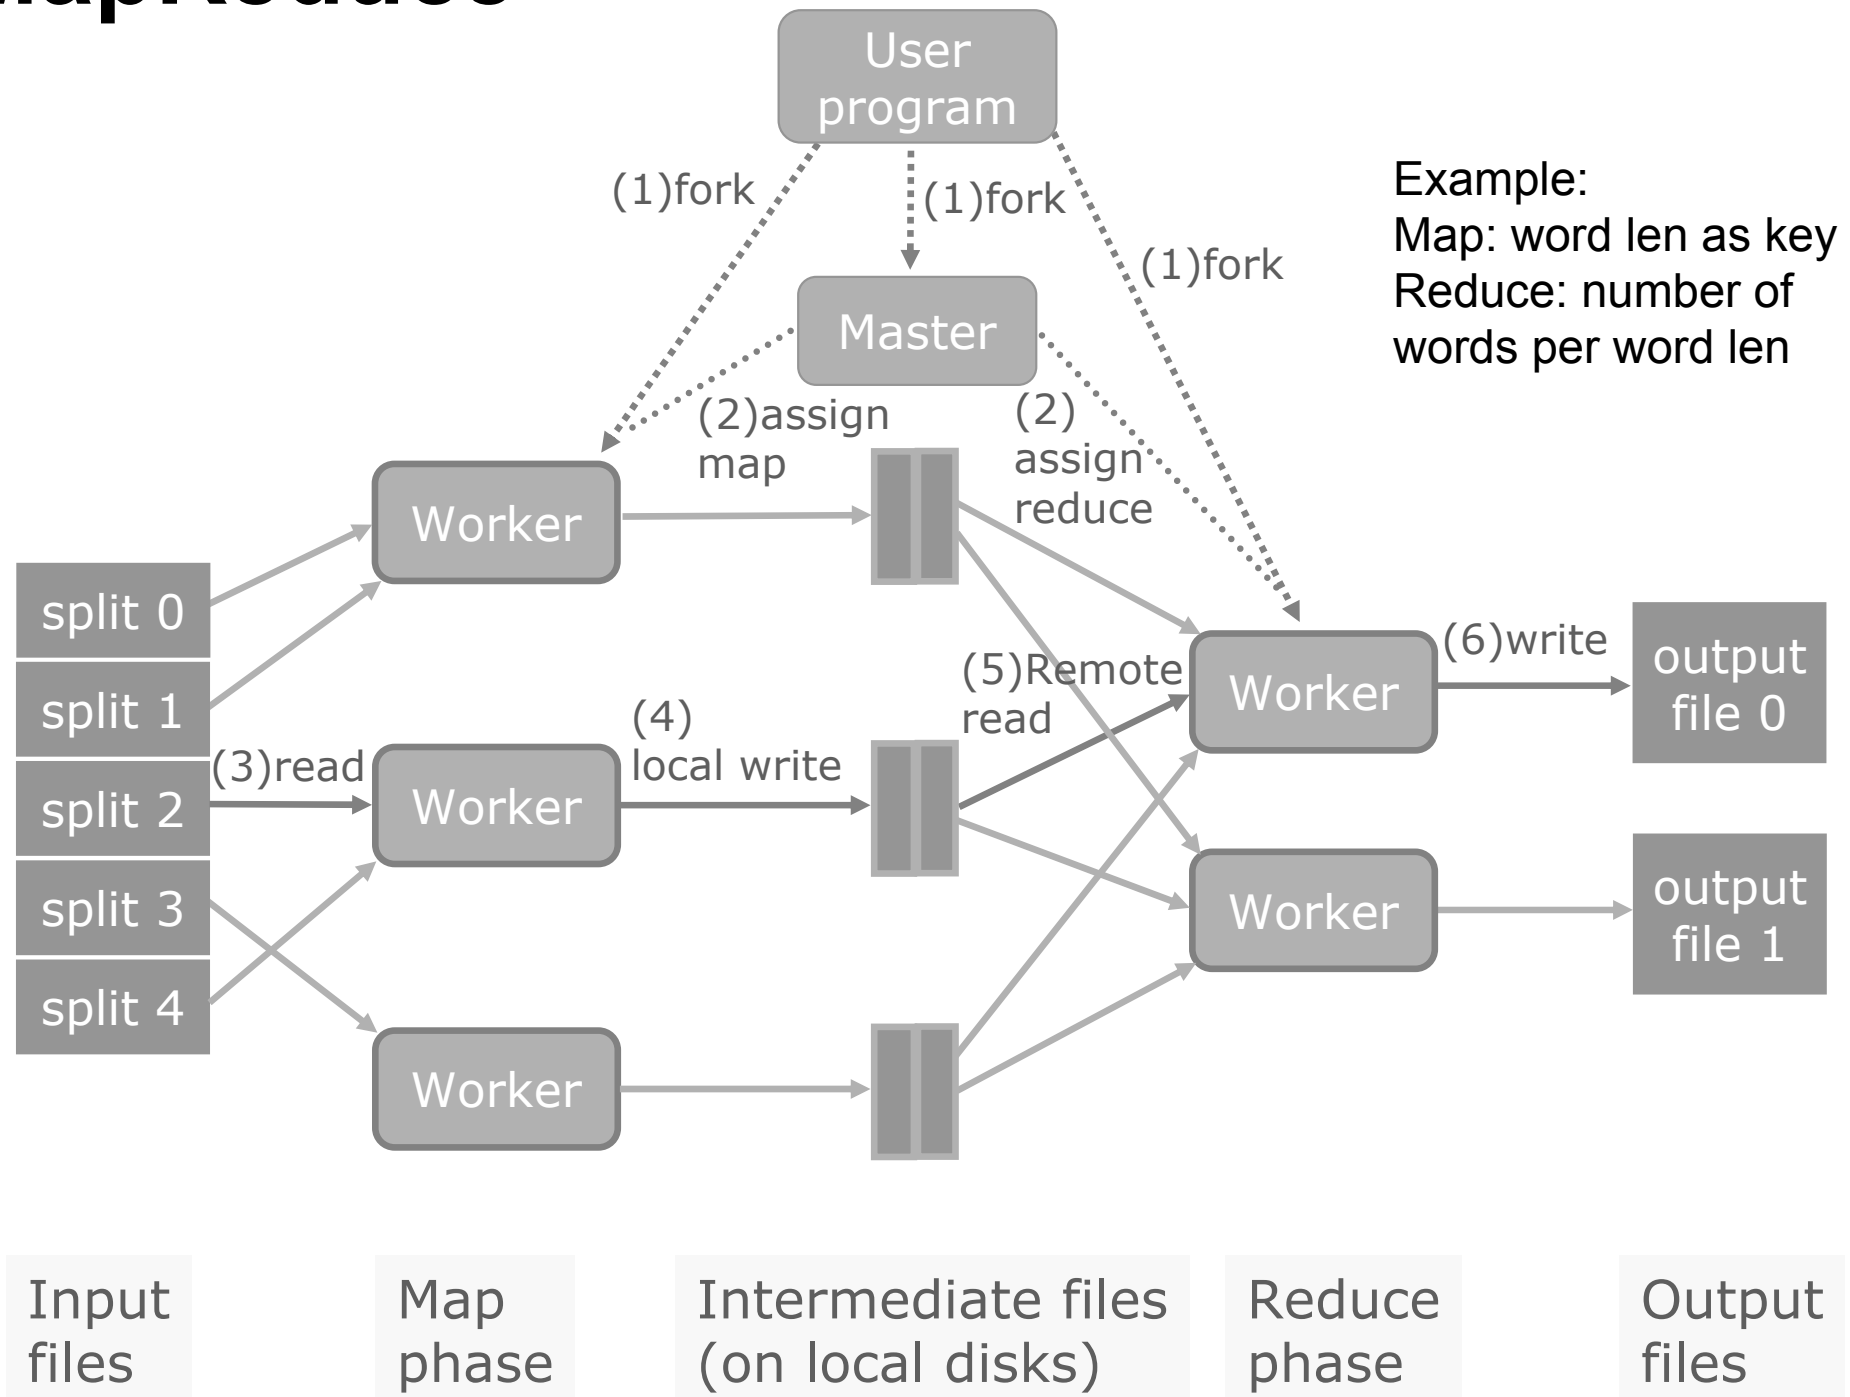
Automatic distribution and parallelization

Fault-tolerance

Cluster management tools

Abstraction for programmers

# MapReduce



# MapReduce terminology

**Job** is a full program that consists of a **Mapper** and **Reducer** for a dataset

**Task** is an execution of a Mapper / Reducer on some data

**Task-in-Progress (TIP)**

**Task Attempt** is an instance of an attempt to run a task on a node

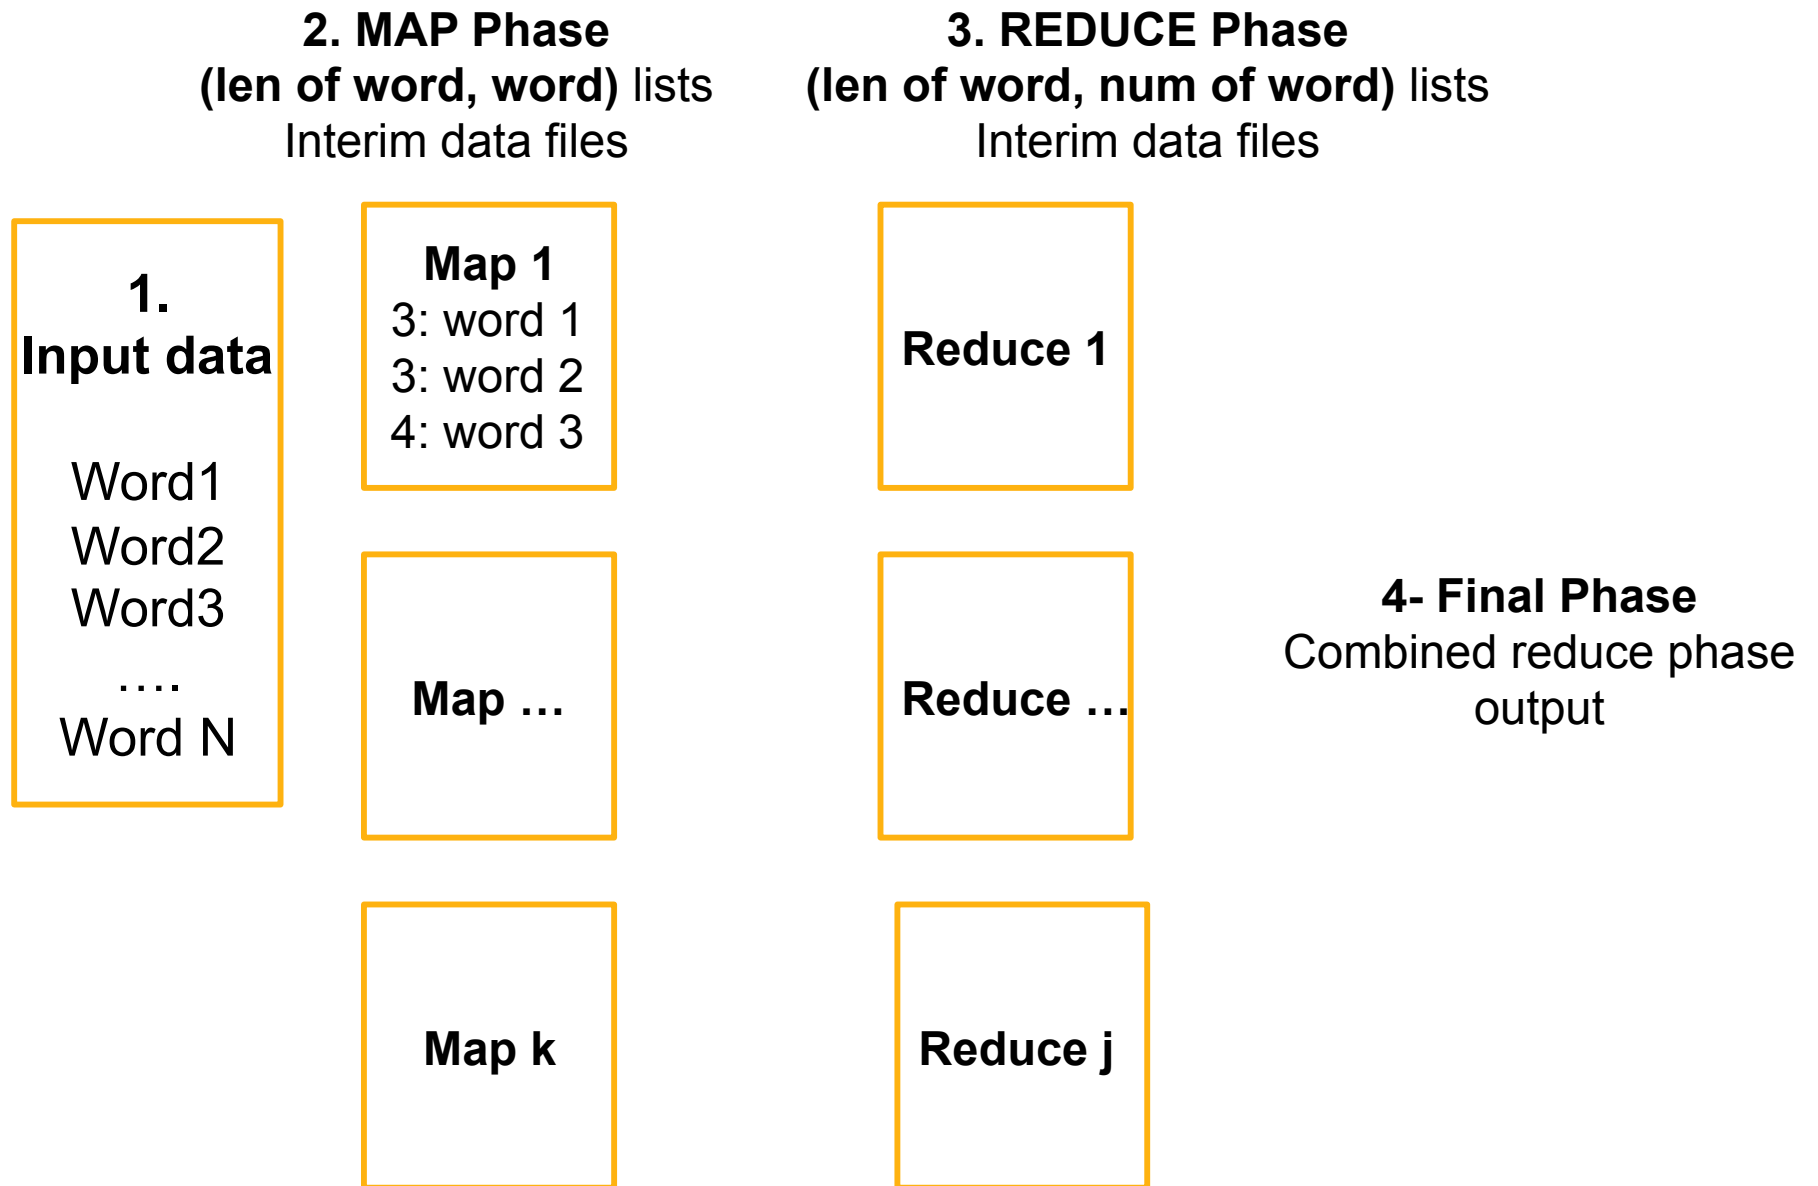
# Task Attempts

A given task is attempted at least once and possibly many times if it crashes

If a task crashes consistently, it will be abandoned eventually

Multiple attempts pertaining to a task may happen in parallel with the **speculative execution** feature

# MapReduce example: counting words per word length





# MapReduce Programming Model


Inspired by functional programming

Two key functions that need to be implemented:

- **map** (in\_key, in\_value) → (out\_key, intermediate\_value) list
  - Data source records are fed as key,value pairs
  - map() produces one or more intermediate values with an output key
- **reduce** (out\_key, intermediate\_value list) → out\_value list
  - Intermediate values for given key are combined into a list
  - Reduce() combines those values into one or more final values for the same output key
  - Optional and not needed by all applications


# map

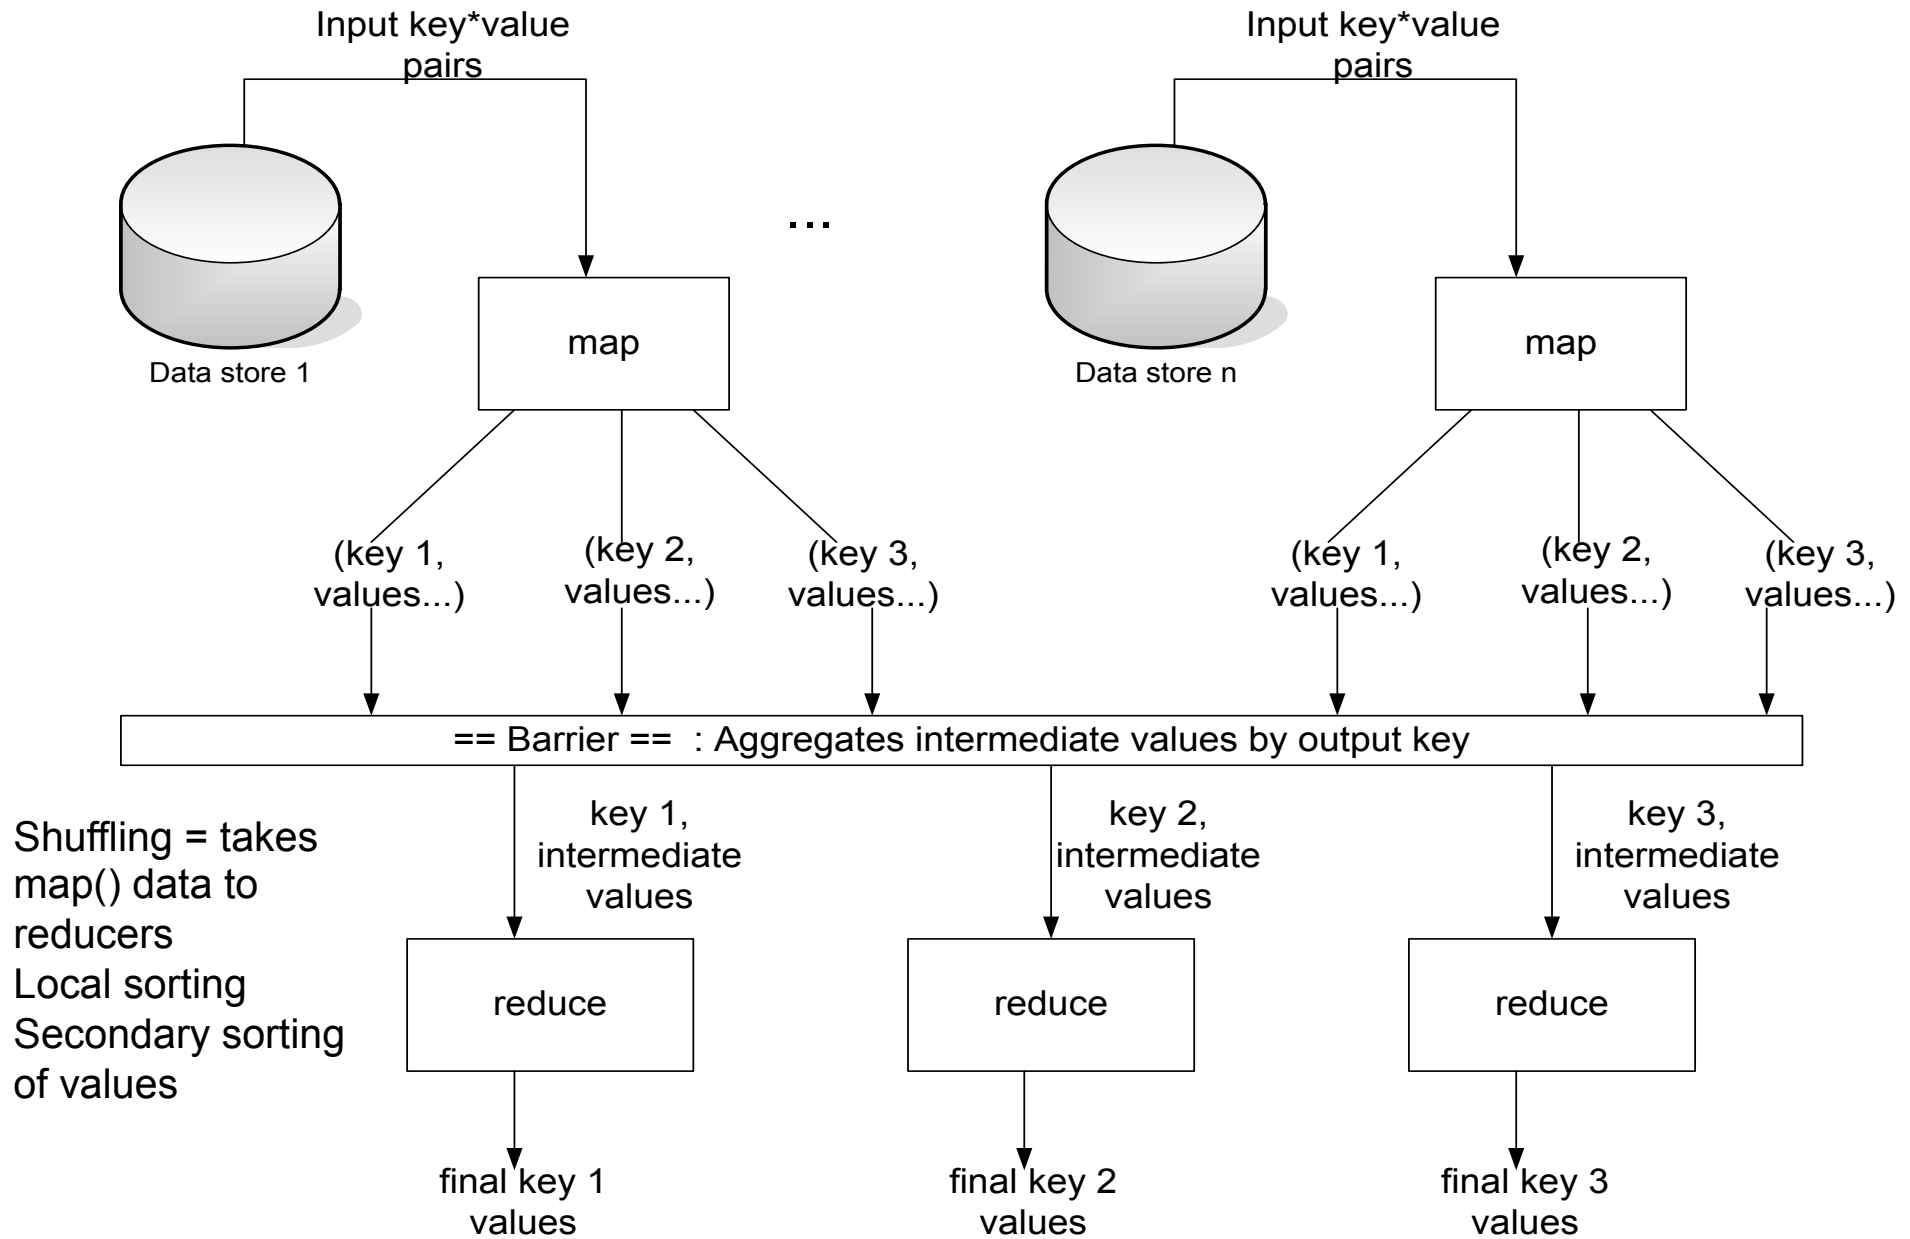
```
map (in_key, in_value) ->  
    (out_key, intermediate_value) list
```



# Reduce

```
reduce (out_key, intermediate_value list) ->  
      out_value list
```





Source: <https://courses.cs.washington.edu/courses/cse490h/08au/lectures/mapred.pdf>

# Reduce Details

- Three key phases
  1. Reducer copies sorted output (based on key) from each Mapper using HTTP
  2. The framework sorts Reducer inputs by keys, because different Mappers have emitted the same key. The shuffle and sort phases happen simultaneously. SecondarySort can be used to sort values.
  3. Reduce is applied for each key in the sorted inputs

Reducer output is not sorted.

# MapReduce Parallelism

- Both map() and reduce() operations are executed in parallel
  - Map() creates intermediate values from input
  - Reduce() functions operate on different output keys
  - Values are independently processed
- 
- Note that reduce cannot start before map is finished
  - Secondary sort on values: the application can extend key with a secondary key and define a grouping comparator

# Locality

The map() task input data is divided into GFS/HDFS blocks

Master program distributes tasks based on the location of data

Map() tasks should be run on the same machine as the physical data file (or the same rack)

# Fault Tolerance

Fault tolerance is realized by periodic heartbeats

Master pings worker nodes to detect node failures

Master must re-execute tasks that have failed

Master can notices if particular input key/values cause problems in map and can skip those



# Node failures

In the worst-case scenario the master compute node fails causing the job to be restarted

Other failures can be managed by the master

Failure of a worker node requires that the task is assigned to another worker

Master reschedules failed tasks when workers become available

# Optimizations

Reduce tasks cannot start before the whole map phase is complete

Thus single slow machine can slow down the whole process

Master can execute many redundant map tasks and then use the results of the first task to complete

# Optimizations: Combining Phase

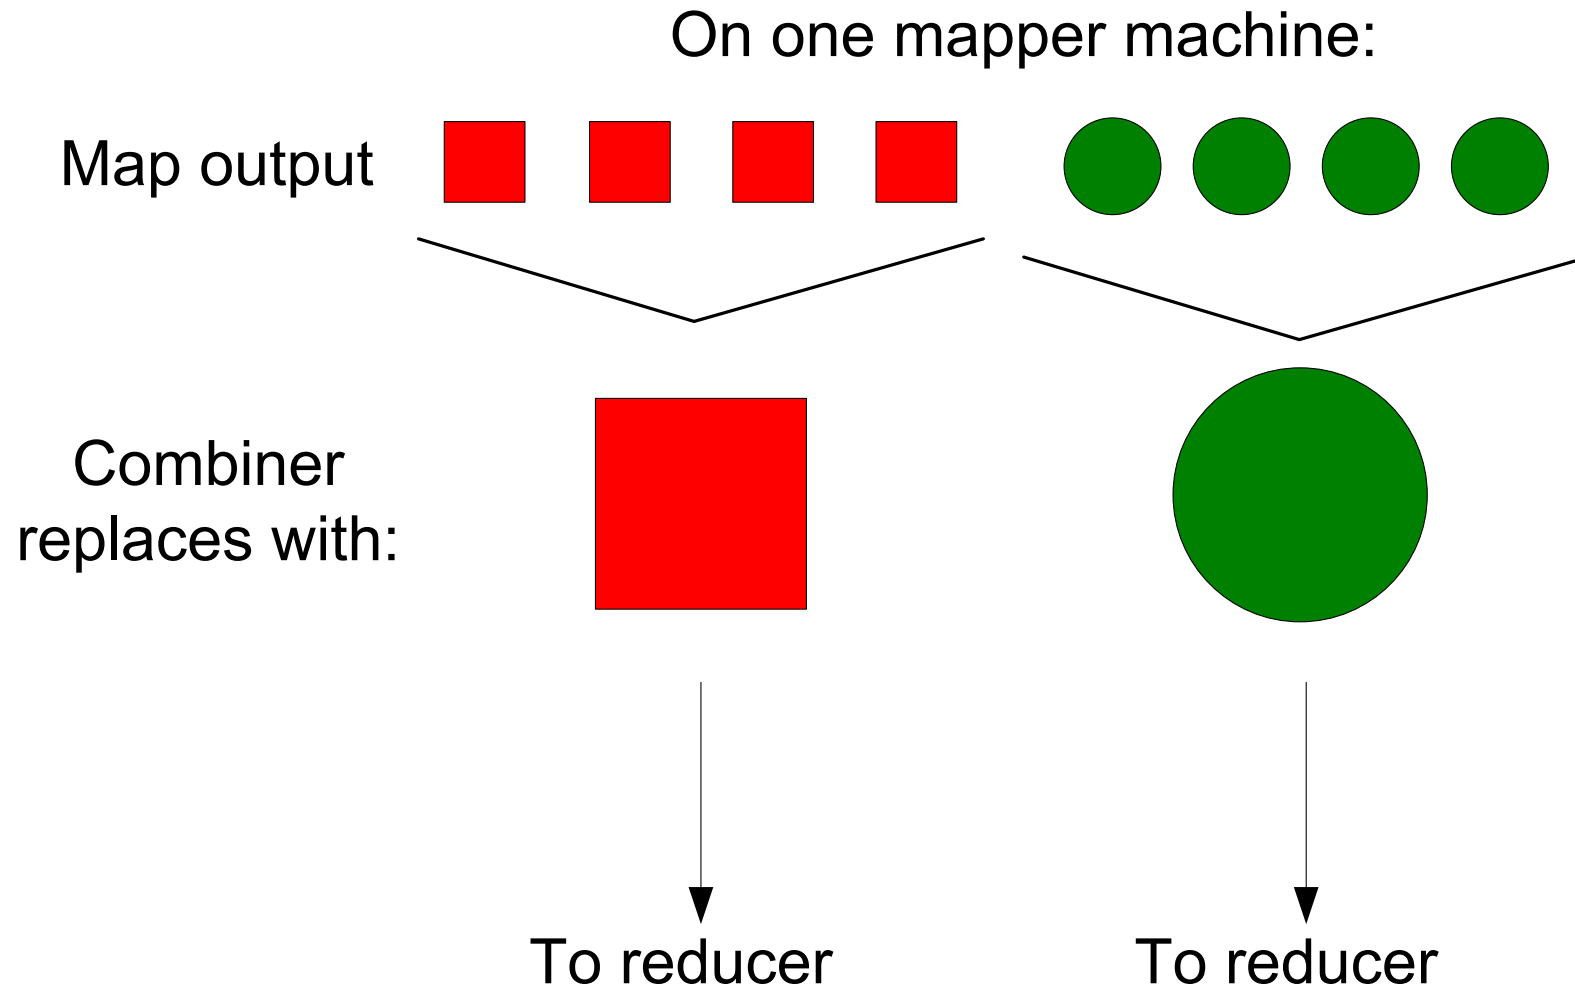
Performance can be increased by running a mini reduce phase on local map output

Executed on mapper nodes after map phase

Saves bandwidth before sending data to a full reducer

Reducer can be a combiner if it is commutative and associative

# Combiner, graphically



Source: <https://courses.cs.washington.edu/courses/cse490h/08au/lectures/mapred.pdf>

# Partitioner

Partitioner divides the **intermediate key space**

Assigns intermediate key-value pairs to reducers

Thus **n partitions results in n reducers**

Between map and reduce phases:

- data is shuffled: parallel-sorted and exchanged

- data is moved to the correct shard for reducing

- partition function accepts the key and the number of reducers and then returns the index of the reducer

Supports load balancing

# MapReduce Summary

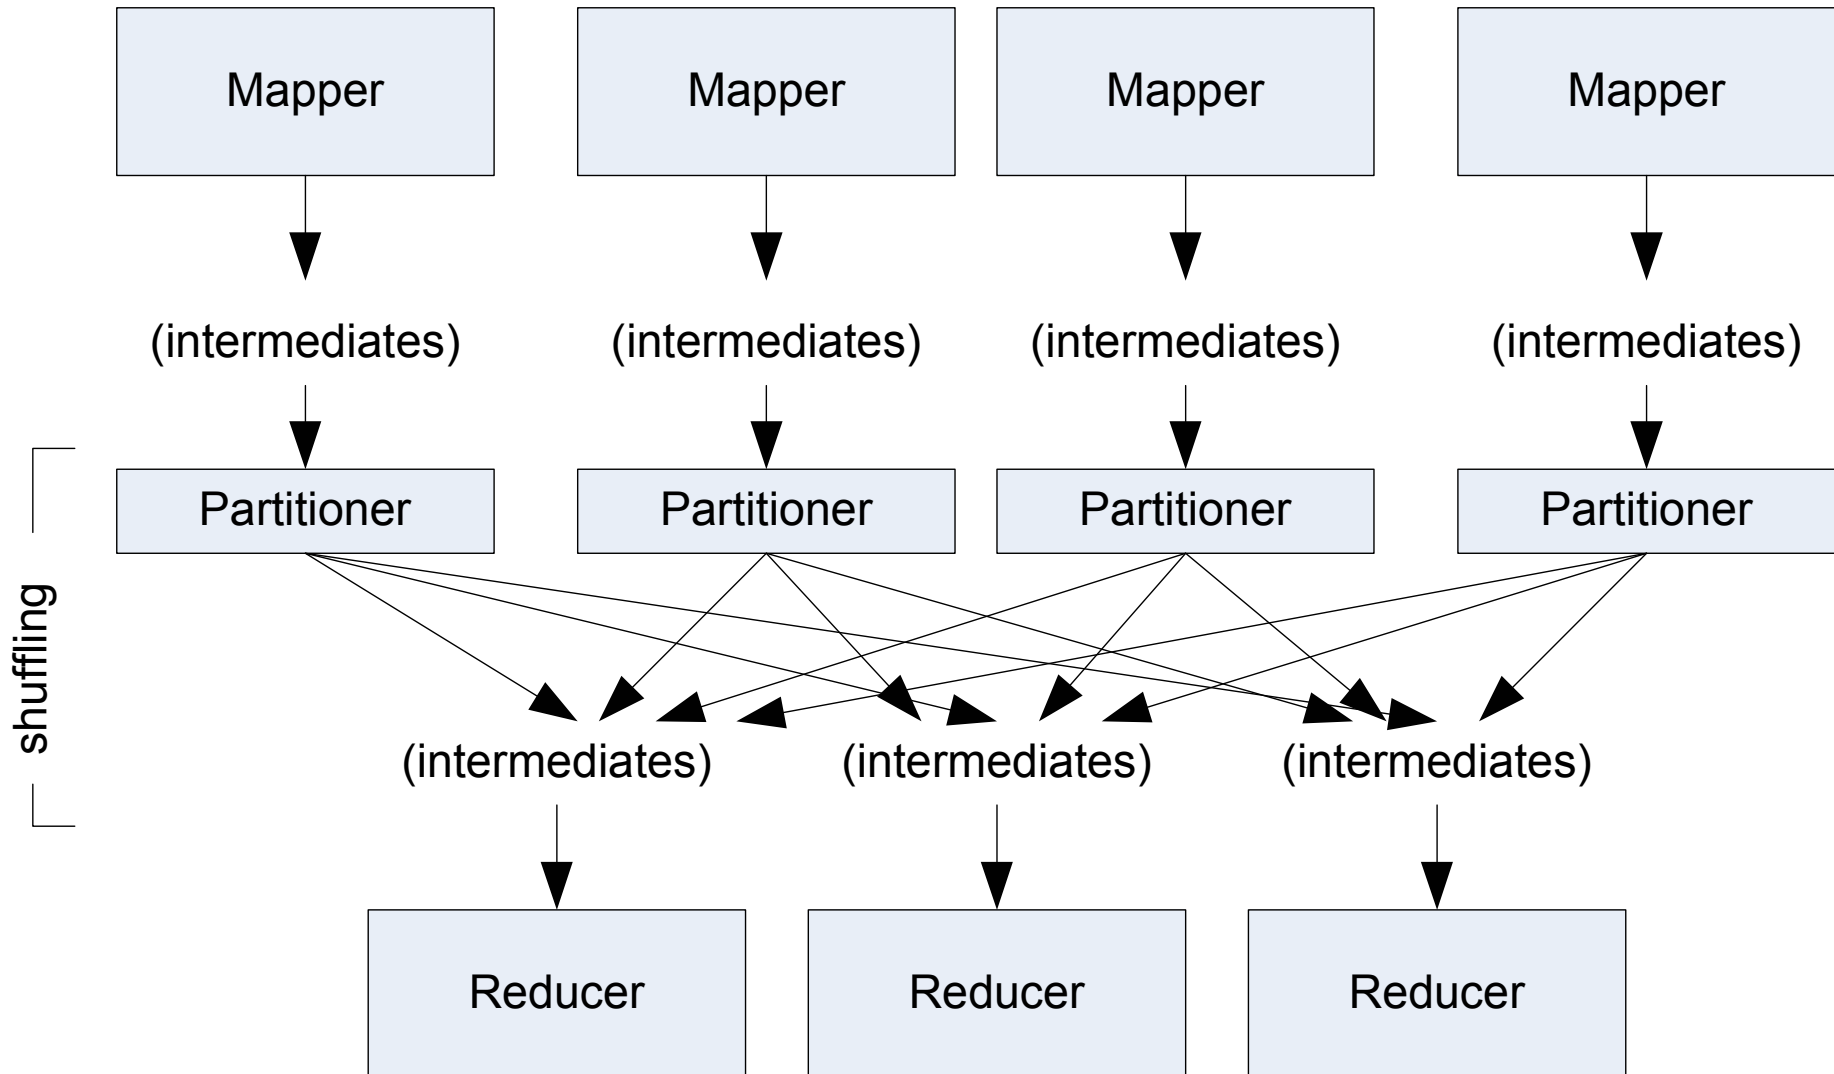
Two key functions that need to be implemented:

- **map** (in\_key, in\_value) → (out\_key, intermediate\_value) list
- **reduce** (out\_key, intermediate\_value list) → out\_value list

With two optimizations:

- **combine** (key, intermediate\_value list) → intermediate\_out\_value list
- **partition** (key, number of partitions) → partition for key

# Partition and Shuffle



# Data management as part of MapReduce

## Purlieus and CAM

Data management as part of MapReduce

Need to know task types (map/reduce intensive)

Map intensive jobs benefit from locality awareness

Up to 80% reduction in execution time



# Key algorithms for MapReduce

Inverted Index

Sorting

PageRank

Sorting

Searching

Statistics

Average, SD, count

Advanced algorithms

# Filtering Algorithms

Finding files or items with specific characteristics

Searching for patterns in web logs or files

Filtering is mostly done in the map phase

Reduce can be simply the identity

# Aggregation Algorithms

Computing the minimum, maximum, sum, average, ...  
of the given values

Count the number of Tweets per day

Map is simple or the identity, reduce is doing most of  
the work

# K-Means Clustering Algorithm

Iterative algorithm that is run until it converges

1. K initial points (centers) are chosen at random.
2. K clusters are formed by associating every data point (observation) with the nearest center.
3. For each cluster, recompute the centers (determine centroid)
4. Repeat from 2 until convergence.

Source: Riccardo Torlone. Analytics on Big Data.  
Universita Roma Tre.

# K-Means for MapReduce

## Map phase

Each map reads the K centroids and a block from the input dataset

Each point is assigned to the closest centroid

Output: <centroid, point>

## Reduce phase

Obtain all points for a given centroid

Recompute the new centroid

Output: <new centroid>

## Iteration:

Compare the old and new set of K centroids

If they are similar then Stop

Else Start another iteration unless maximum of iterations has been reached.

# Optimizing K-Means for MapReduce

Combiners can be used to optimize the distributed algorithm

Compute for each centroid the local sums of the points

Send to the reducer: <centroid, partial sums>

Use of a single reducer

Data to reducers is very small

Single reducer can tell immediately if the computation has converged

Creation of a single output file

# Limitation: iterative algorithms

MapReduce tasks must be written as acyclic dataflow programs

Stateless mappers and reducers

Batch model

Difficult to implement iterative processing of datasets

Machine learning typically requires iterative operation: the dataset is visited multiple times by the algorithm